



CS 524 A

Introduction to Cloud Computing

Module 3: Virtual Machines (Part 1)

OUTLINE

- Virtual vs. Real
- Computer organization
- OS recap
 - A process as a unit of computation
 - Multiprocessing
 - Virtual views: How a process gets to think that it owns
 - the CPU
 - **infinite** memory
 - all devices to itself

AMERICAN HERITAGE DICTIONARY

DEFINITION


vir·tu·al (vûr□ch□-əl) :

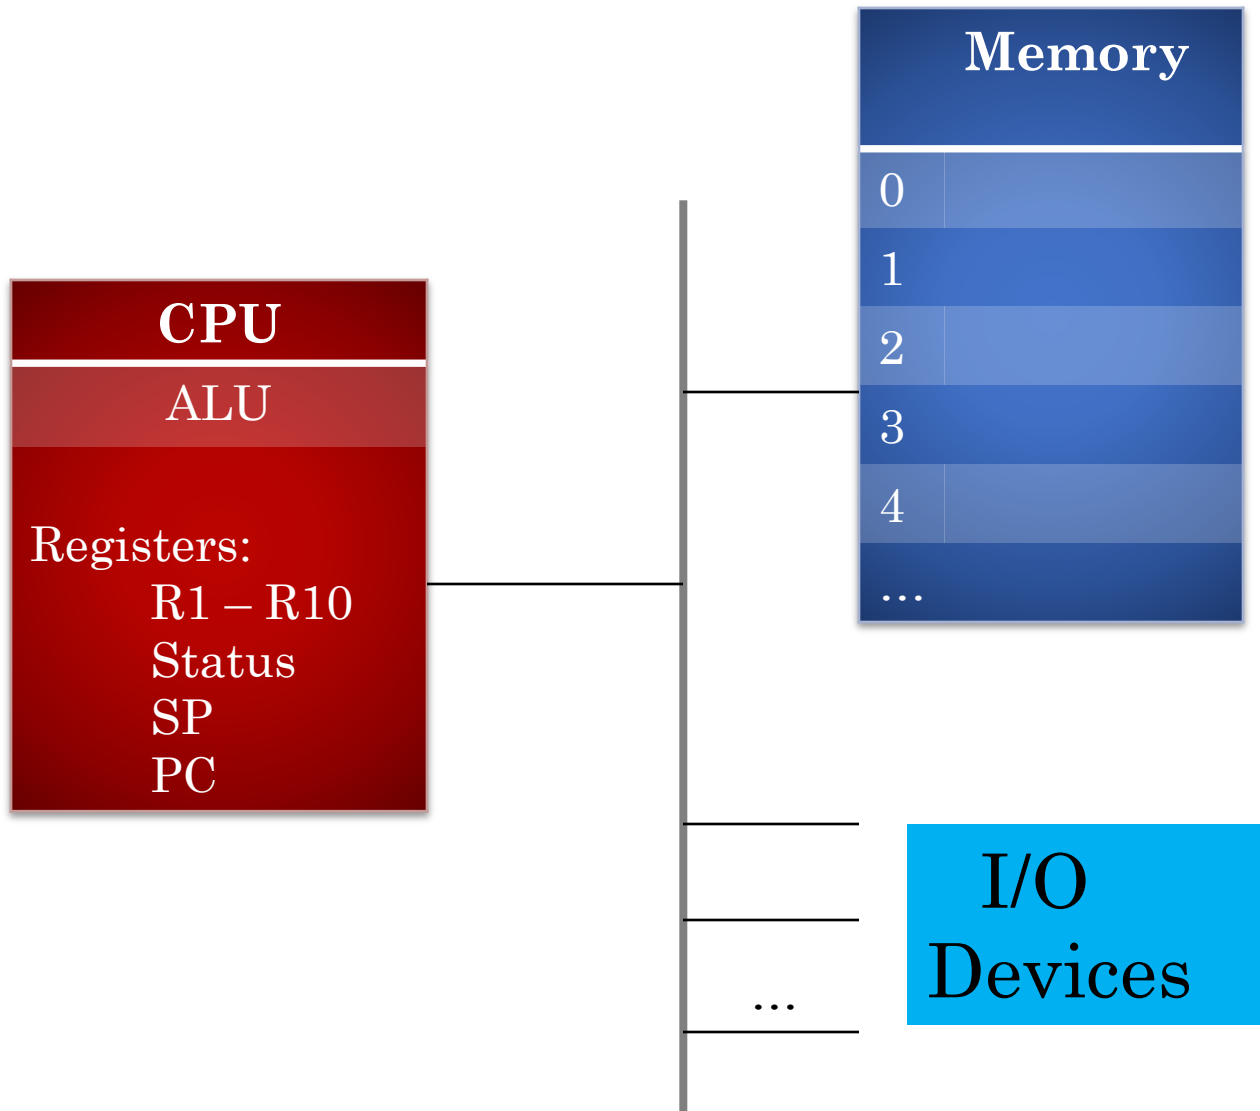
adj. **1.** Existing or resulting in essence or effect though not in actual fact, form, or name: *the virtual extinction of the buffalo.*

2. Existing in the mind, especially as a product of the imagination. Used in literary criticism of a text.

3. *Computers* Created, simulated, or carried on by means of a computer or computer network: *virtual conversations in a chatroom.*

COMPUTER ORGANIZATION





The CPU Loop: first approximation

```
While TRUE
```

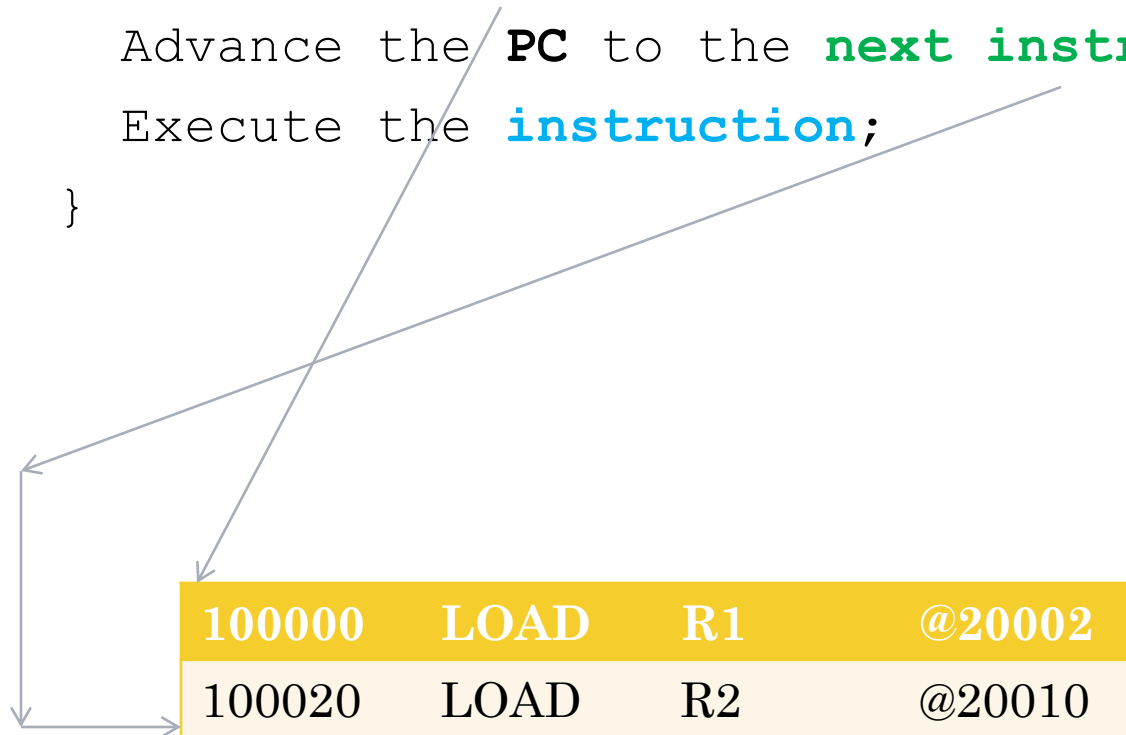
```
{
```

```
    Fetch the instruction pointed to by the PC;
```

```
    Advance the PC to the next instruction;
```

```
    Execute the instruction;
```

```
}
```



A PROCESS

- Modern operating systems support *multiprogramming*—that is an ability to execute several programs concurrently on one CPU or simultaneously on several CPUs
- A *process* is a *program in execution*:
 - A program is a cookbook
 - A CPU is a cook
 - I/O devices are cooking utensils
 - A process is making a dish described in the *program*



Operating System Services To a Process

Program execution (Process
or Thread)

Input/Output (I/O)
operations

File-system support

Interprocess
Communications

Error detection

Resource allocation

Accounting

Protection

The process's stack and the procedure call

Main line of the process code

100000	LOAD	R1	@20002
100010	LOAD	R2	@20010
100020	STORE	R1	@SP
100030	ADD	SP	#-4
100040	STORE	R2	@SP
100050	ADD	SP	#-4
100060	ADD	SP	#-8
100070	JPR	10000000	
100080	ADD	SP	#16
100090

Stack Frame

**Saved PC
(100080)**

Internal variable
2

Internal variable
1

Parameter 2

Parameter 1

The procedure code


10000000	LOAD	R1	@(SP + #20)
...
10005000	RTP		

WHERE IS THIS PROCESS?

- In older systems (up until 1960), there was a place in memory for *one* process. The Operating System loaded it and ran it

Problems:

1. Low CPU Utilization
2. A need to program device drivers in each process for the devices it uses
3. Inability to split a program into manageable independent concurrent pieces
4. Effectively, inability to support more than one interactive user



Multiprogramming

Operating System

Process 1

Process 2

...

Process n

BUT WHILE ONE PROCESS IS WAITING...


- ...Others had better use the CPU and run!



- Then the OS **must** get control of the CPU!

INTERRUPTS AND EXCEPTIONS

1. Interrupts





Interrupts occur *asynchronously* with the execution of the code

INTERRUPTS AND EXCEPTIONS

1 Exceptions (examples)



3) A *TRAP* instruction



1) Memory: Bus error
“*LOAD R1 700000000000*”




2) A non-executable instruction:
“*DIV R1 R2*”

Exceptions occur *synchronously* with the execution of the code

THE OPERATING SYSTEM (KERNEL) CODE

- A library of subroutine calls, which are performed by the **CPU itself**
- Addresses of these service routines reside in the interrupt vector table, indexed by the interrupt or exception number.

	Interrupt Vector Table
0	Disc service routine
1	Keyboard service routine
2	Div by 0 exception routine
...	



An example: Using a TRAP instruction to set a breakpoint

ADD R1, #1	2A 0001 01
STORE R1, @R2	B3 0001 02
ADD R2, #4	2A 0004 02

Save the instruction and replace it with *TRAP #1*

ADD R1, #1	2A 0001 01
TRAP #1	F1 0001 02
ADD R2, #4	2A 0004 02



The second approximation of the CPU loop

While TRUE

```
{  
    Fetch the instruction pointed to by the PC;  
    Advance the PC to the next instruction;  
    Execute the instruction;  
    If (an exception #x has been raised) OR  
        ((an interrupt #x has been raised) AND interrupts are enabled)  
    {  
        If it is an exception  
            Restore the previous PC value;  
        Save the STATUS Register and the PC @SP;  
        PC = Interrupt_Vector[#x];  
    }  
}
```

Interrupt vectors

#1	2300000
#2	7000000
...	
#x	3000000
...	

An interrupt service routine

30000000	DISI
30000004	SAVEREGS @SP
...	...
30006000	RESTREGS @SP
30006004	ENI
30006008	RTI

An interrupt stack frame

...
Saved PC
Saved Status Register
Saved registers' image

A simple *hypervisor-debugger*

The *Go()* code

```
Go ()
{
    #DISI; /* disable interrupts */
    registers_struct.SP =
        registers_struct.SP +4;

    #RESTREGS @registers_struct;
    #STORE PC    @SP+8;
    #STORE STATUS @SP+4
    /* place the PC and Status in
       their proper place within
       the stack frame */

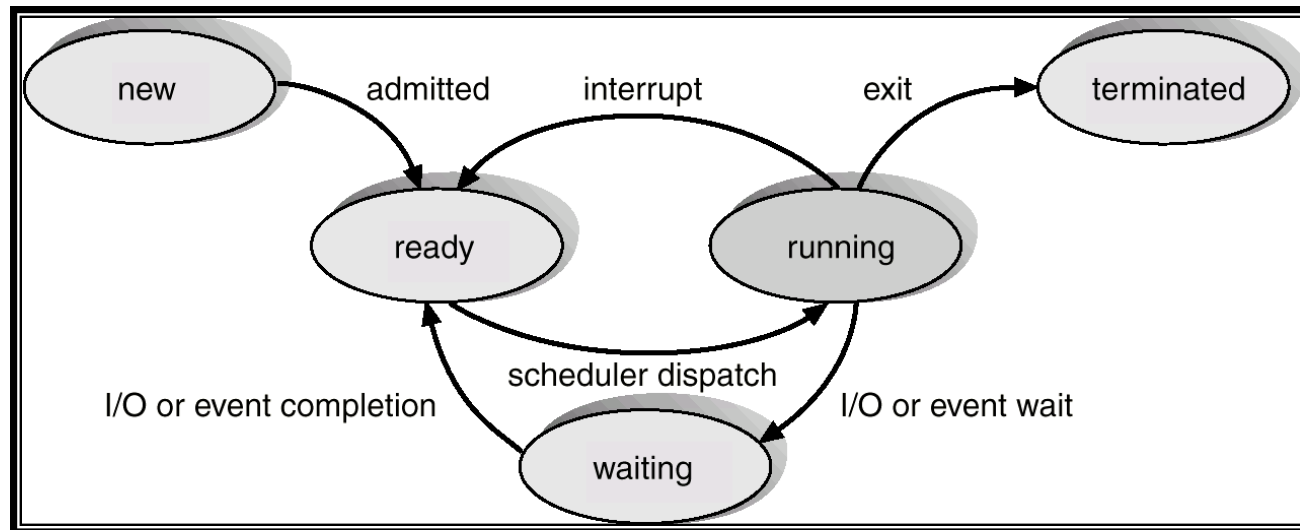
    #ENI; /* enable interrupts */
    #RTI;
}
```

The service routine for *TRAP #1*

The TRAP #1 vector stores the address of the TRAP_1_Service Routine:

```
#DISI; /*disable interrupts*/
#SAVEREGS @registers_struct;
#ENI; /*enable interrupts */
display(registers_struct, PC,
        STATUS);
command_line();
```

The process lives through these states

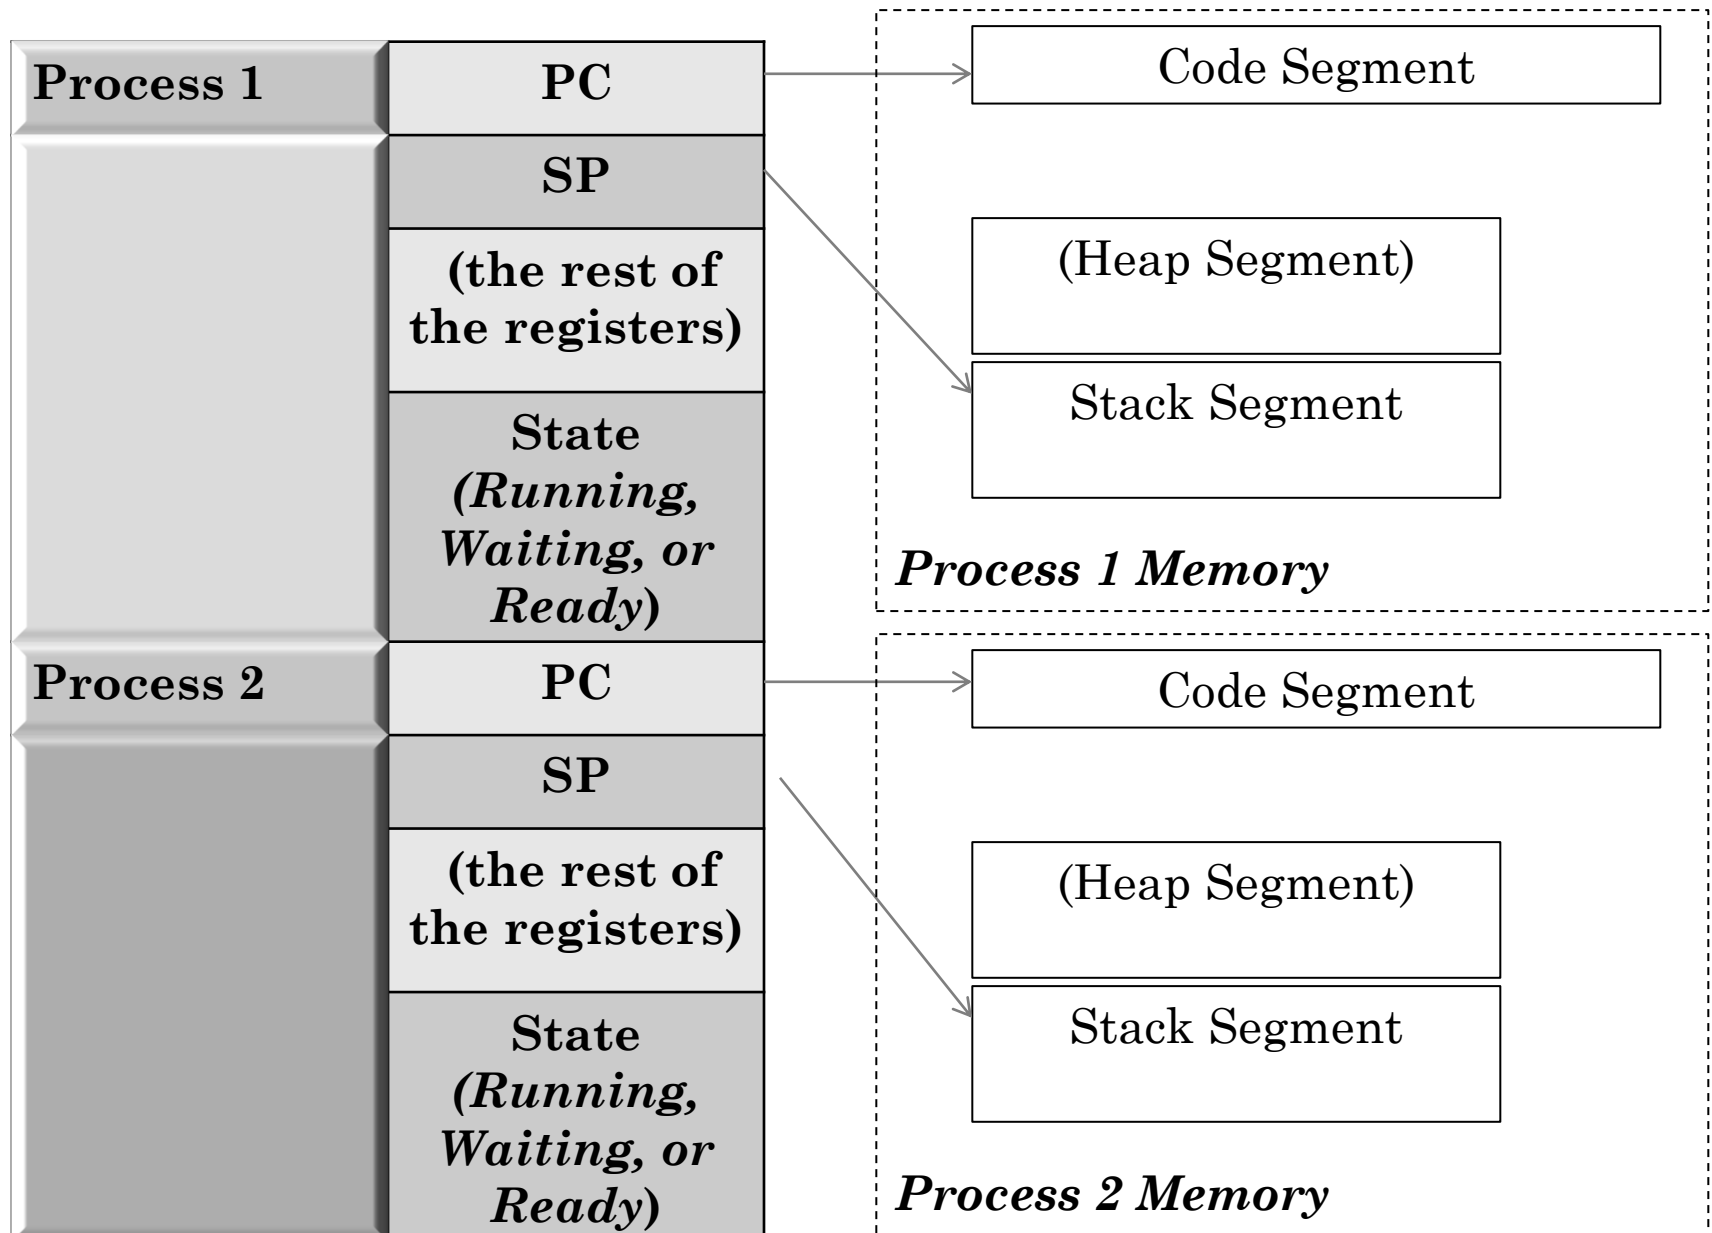


From Silberschatz & al.:
Operating Systems Concepts

THE PROCESS INFORMATION

Process ID
Process state (<i>ready, running, etc.</i>)
All the registers (including PC and SP)
Memory management information
File information
Children processes information
... many other things

The process table




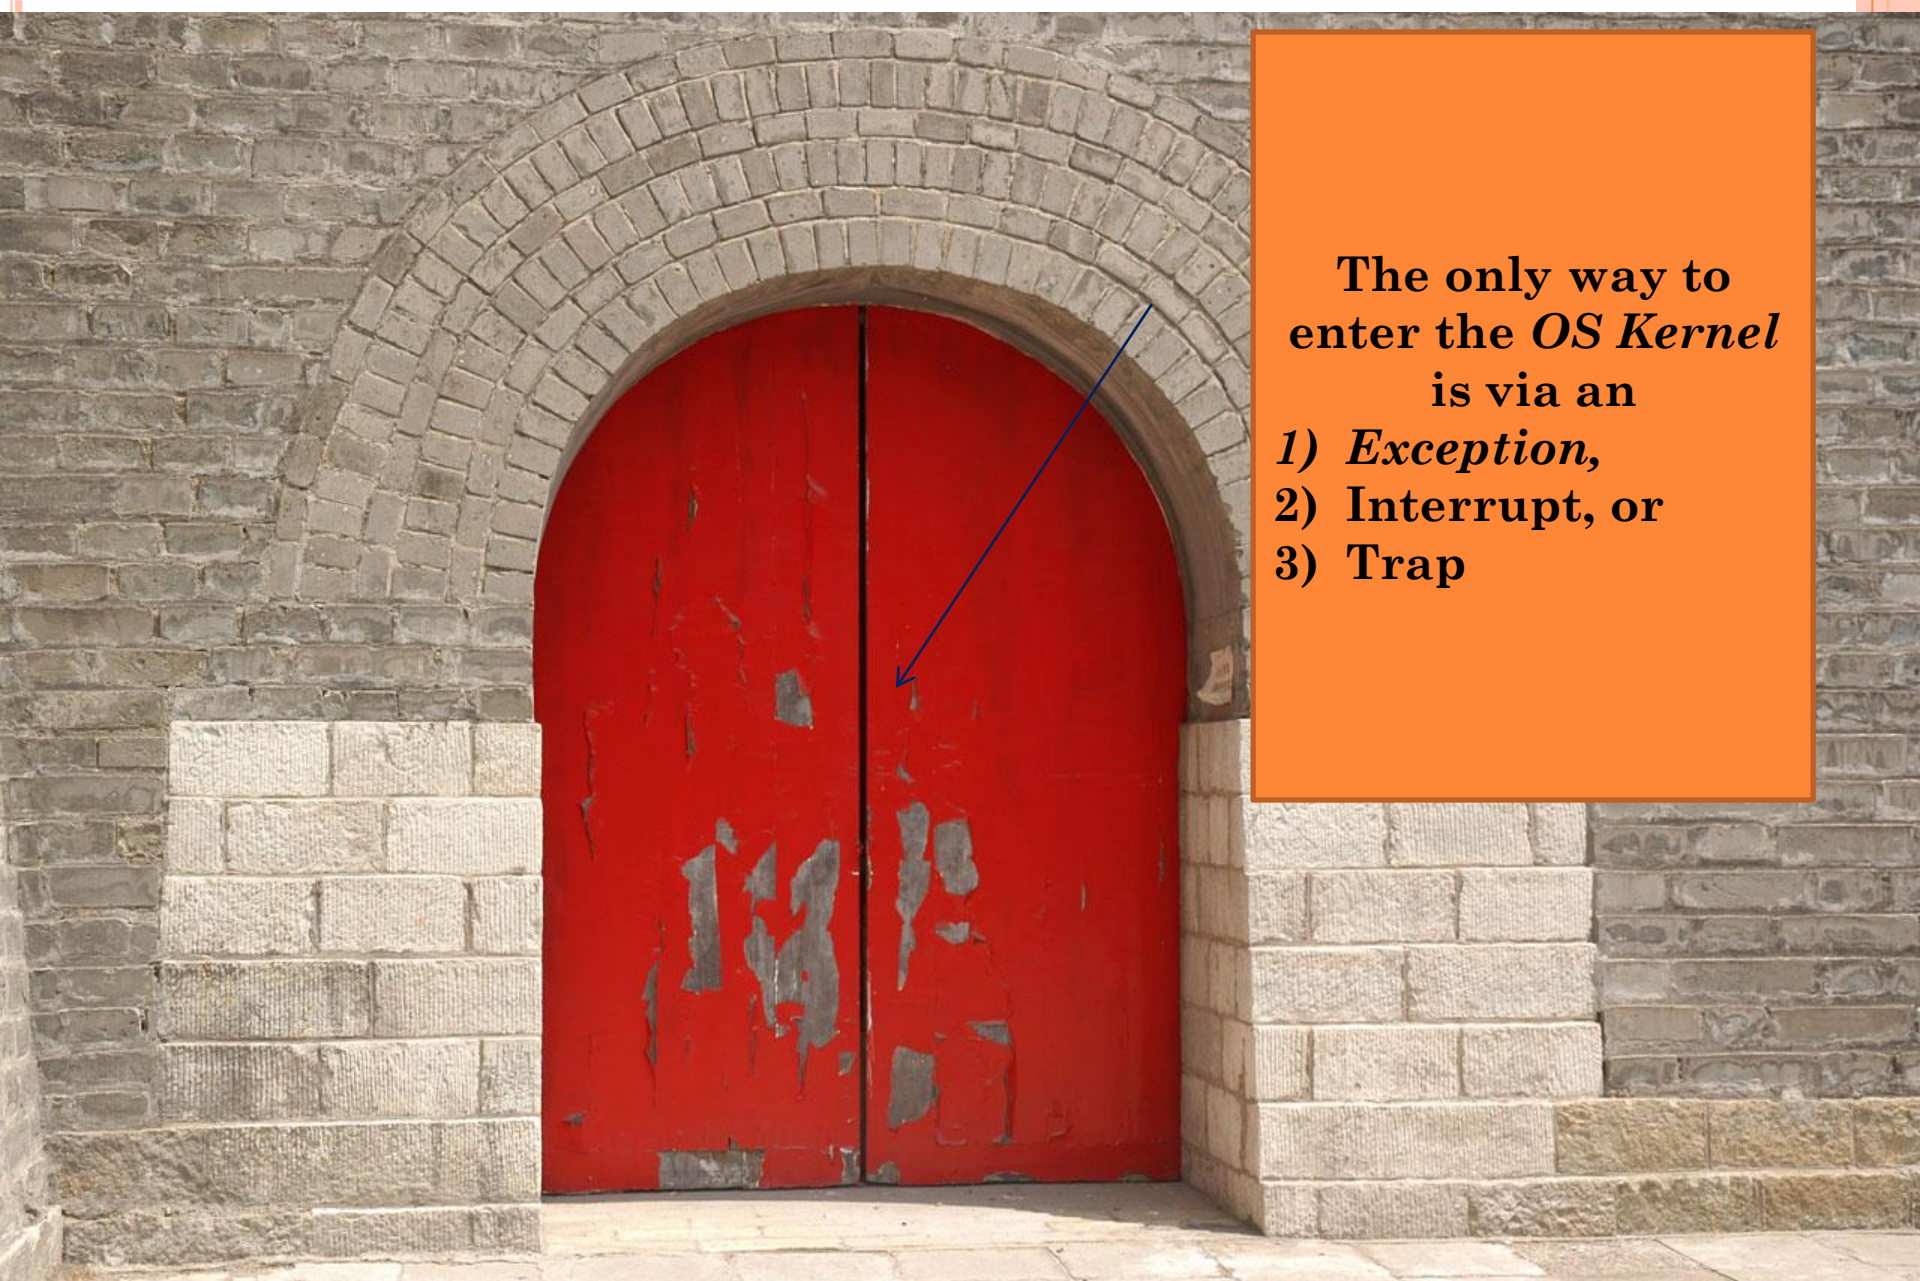
PRIVILEGED INSTRUCTIONS

- Instructions that deal with processing interrupts, changing the status register, performing memory management, and the like are mission-critical.
- Critical instructions, by a long-established convention, require CPU processing in a special—***supervisory*** or ***system***—mode
- A CPU may also have a special set of registers, reserved for the system mode. A separate, **supervisory stack pointer** points to a separate stack
- All exception processing is performed **only** in supervisory mode; an attempt to execute a **privileged instruction in a user mode** is either **ignored [BAD!]** or it **forces an exception**
- (A *Unix* example) Each process therefore actually has two stacks: a ***user stack*** and a ***system stack***.

CPU MODE (AN ESSENTIAL SECURITY FEATURE)

- To ensure that *only* the OS can execute system code (interrupt processing, memory manipulation, etc.), modern CPU execute the system and user code in different modes
- The mode is typically indicated by a flag in the STATUS register

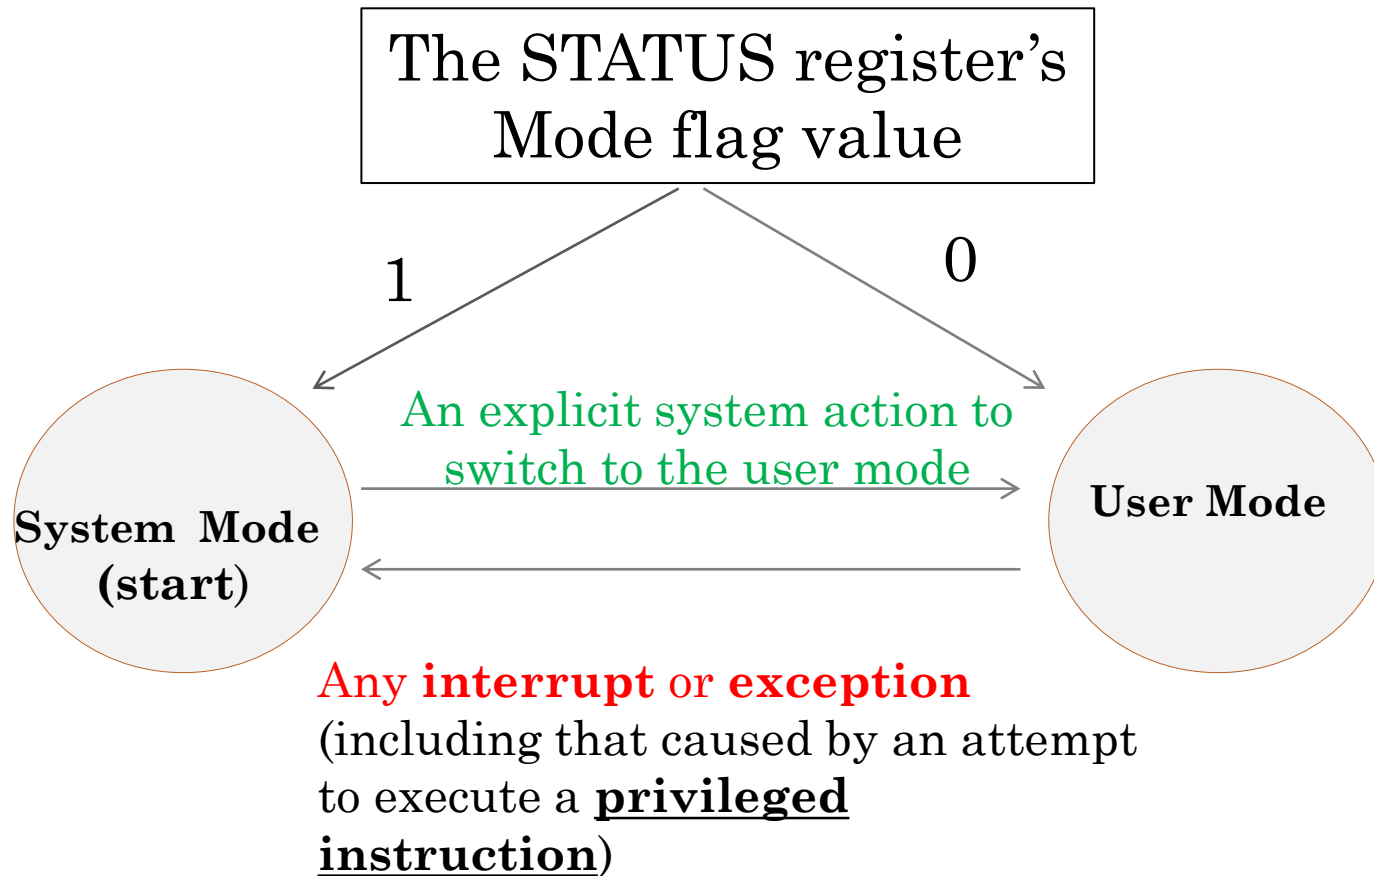


A photograph of a red double door set within a stone archway. The door is made of wood and has a vertical seam down the middle. The paint on the door is peeling and chipped, revealing the underlying wood. A blue arrow points from the text box on the right towards the door. The archway is constructed from grey stone blocks, and the surrounding wall is also made of stone.

**The only way to
enter the *OS Kernel*
is via an**

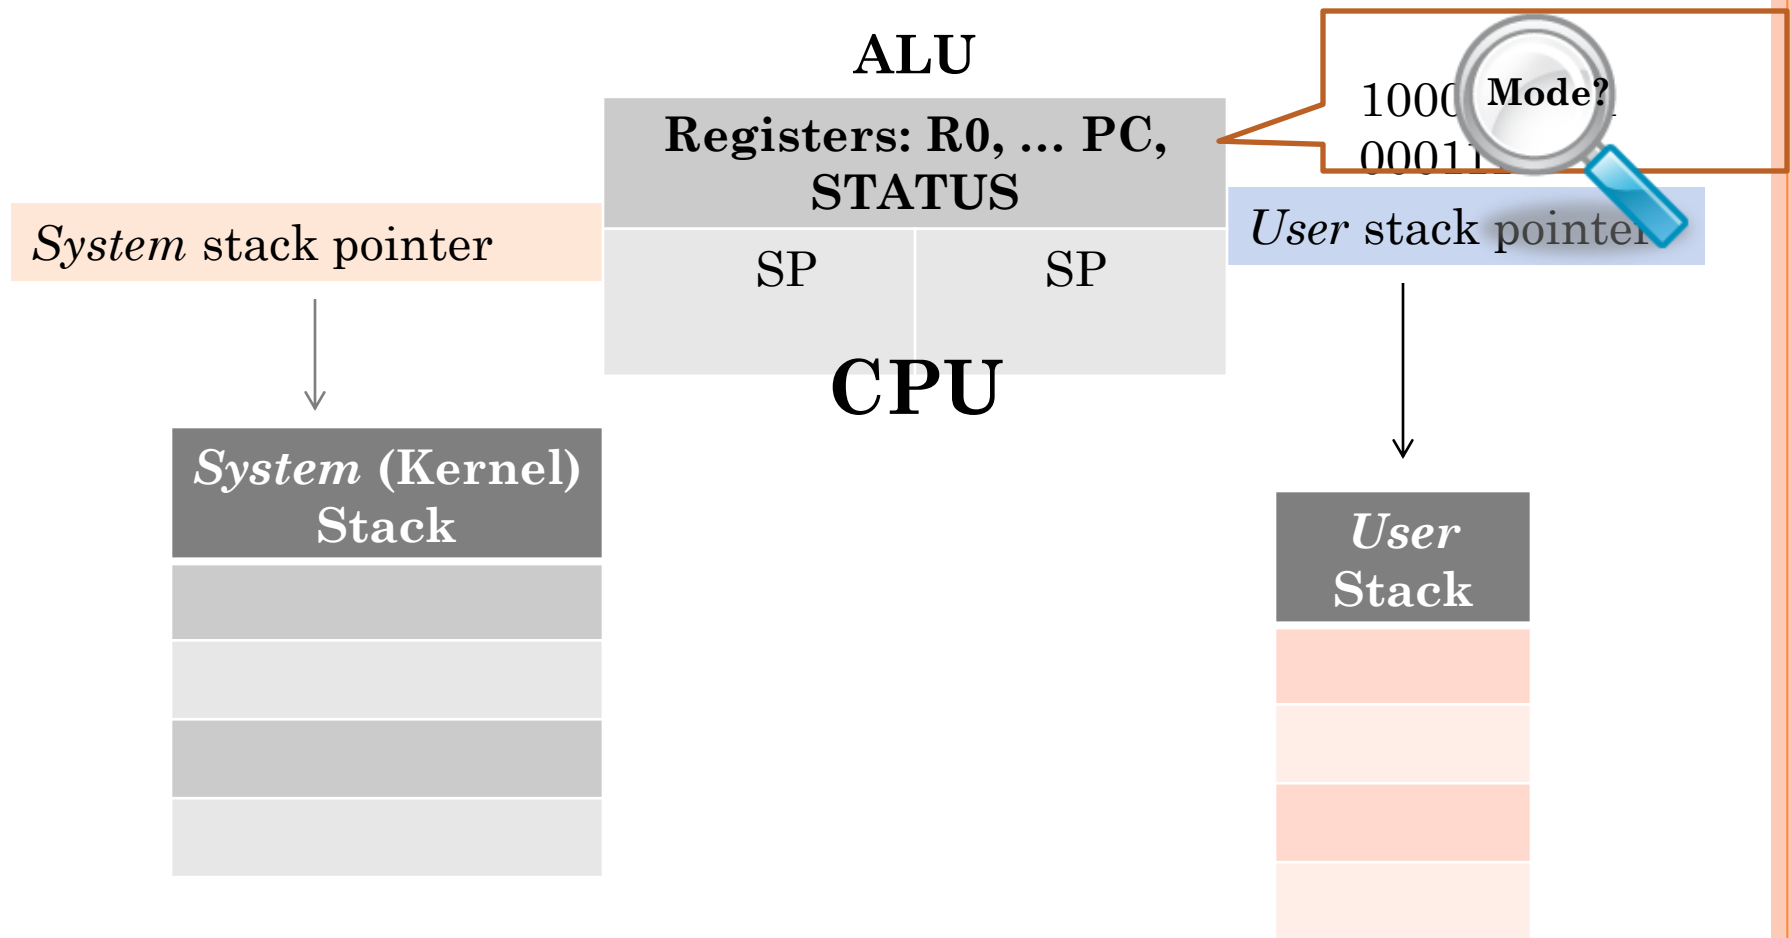
- 1) *Exception*,**
- 2) Interrupt, or**
- 3) Trap**

The *CPU mode* state machine



The System Mode and the User Mode are associated with separate stacks

The modified CPU and the two process stacks



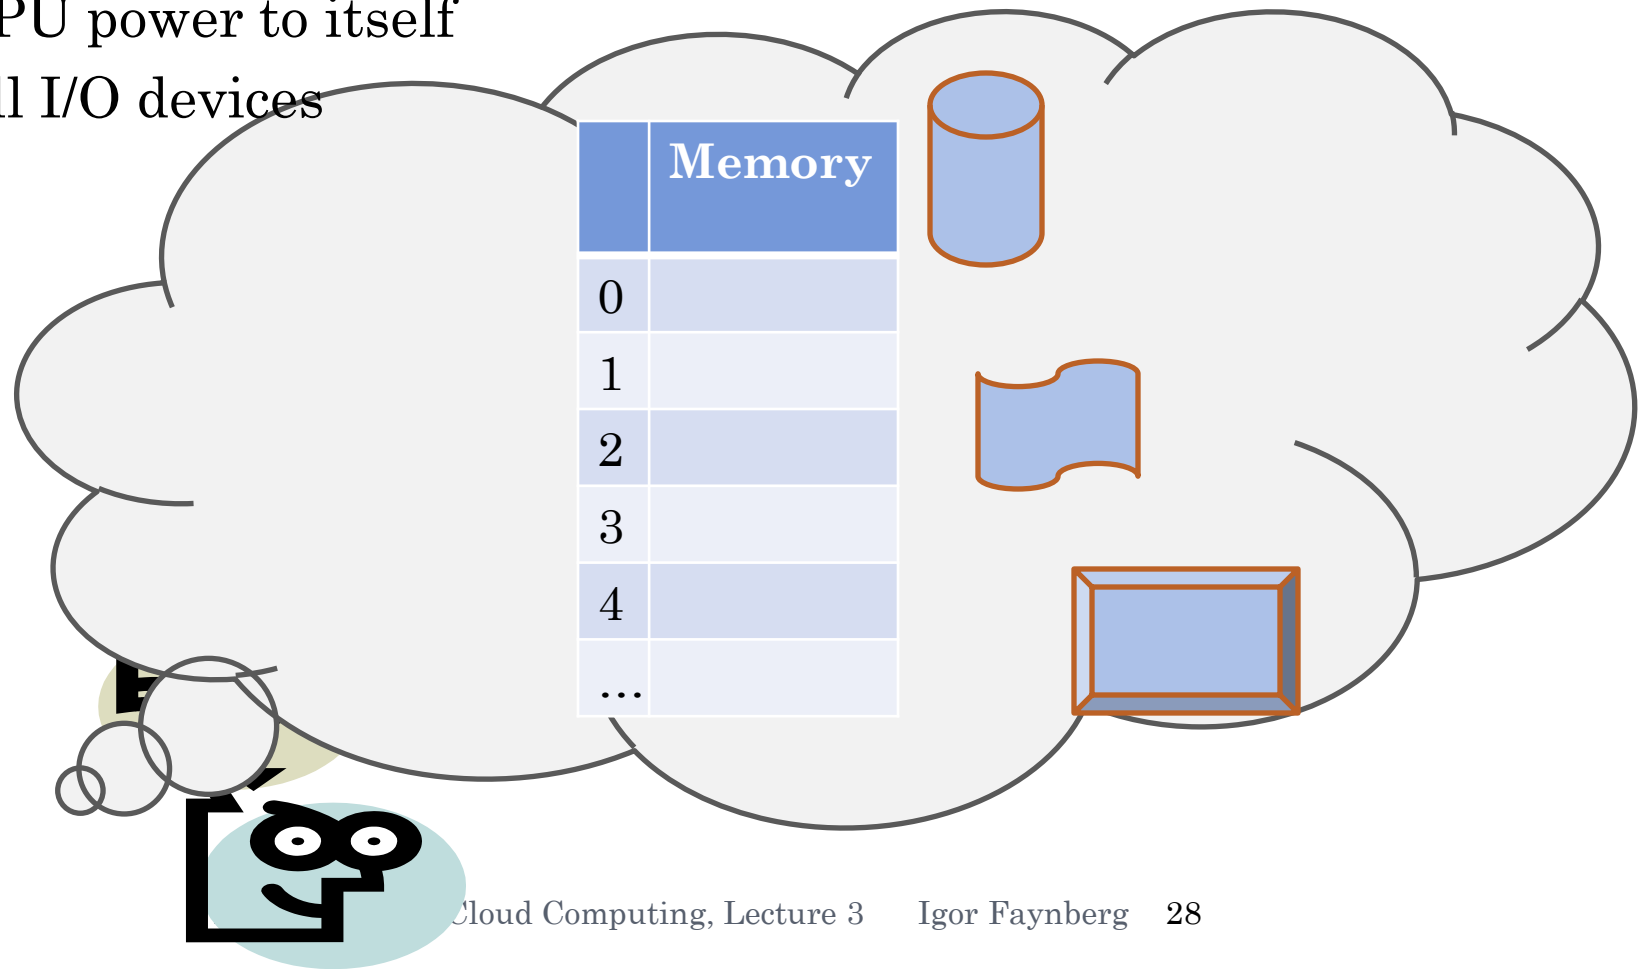
The CPU loop—the final version

```
While TRUE
{
    Fetch an instruction pointed to by the PC;


    If the instruction is valid AND
       the instruction is appropriate for the present mode AND
       the parameters are valid for the operation
    {
        Advance the PC to the next instruction;
        Execute the instruction;
    }
    else
        raise an appropriate exception;
    If (an exception #x has been raised) OR
       (an interrupt #x has been raised) AND interrupts are enabled
    {
        Save the STATUS register and PC on the system stack (@SP);
        Switch to the system mode;
        PC = Interrupt_Vector[x];
    }
}
```

VIRTUALIZATION IN MODERN OPERATING SYSTEMS


- Each *process* thinks that it has
 - huge contiguous memory, starting from address 0
 - CPU power to itself
 - All I/O devices



HOW IS THIS MEMORY *VIRTUALIZATION* ACHIEVED?



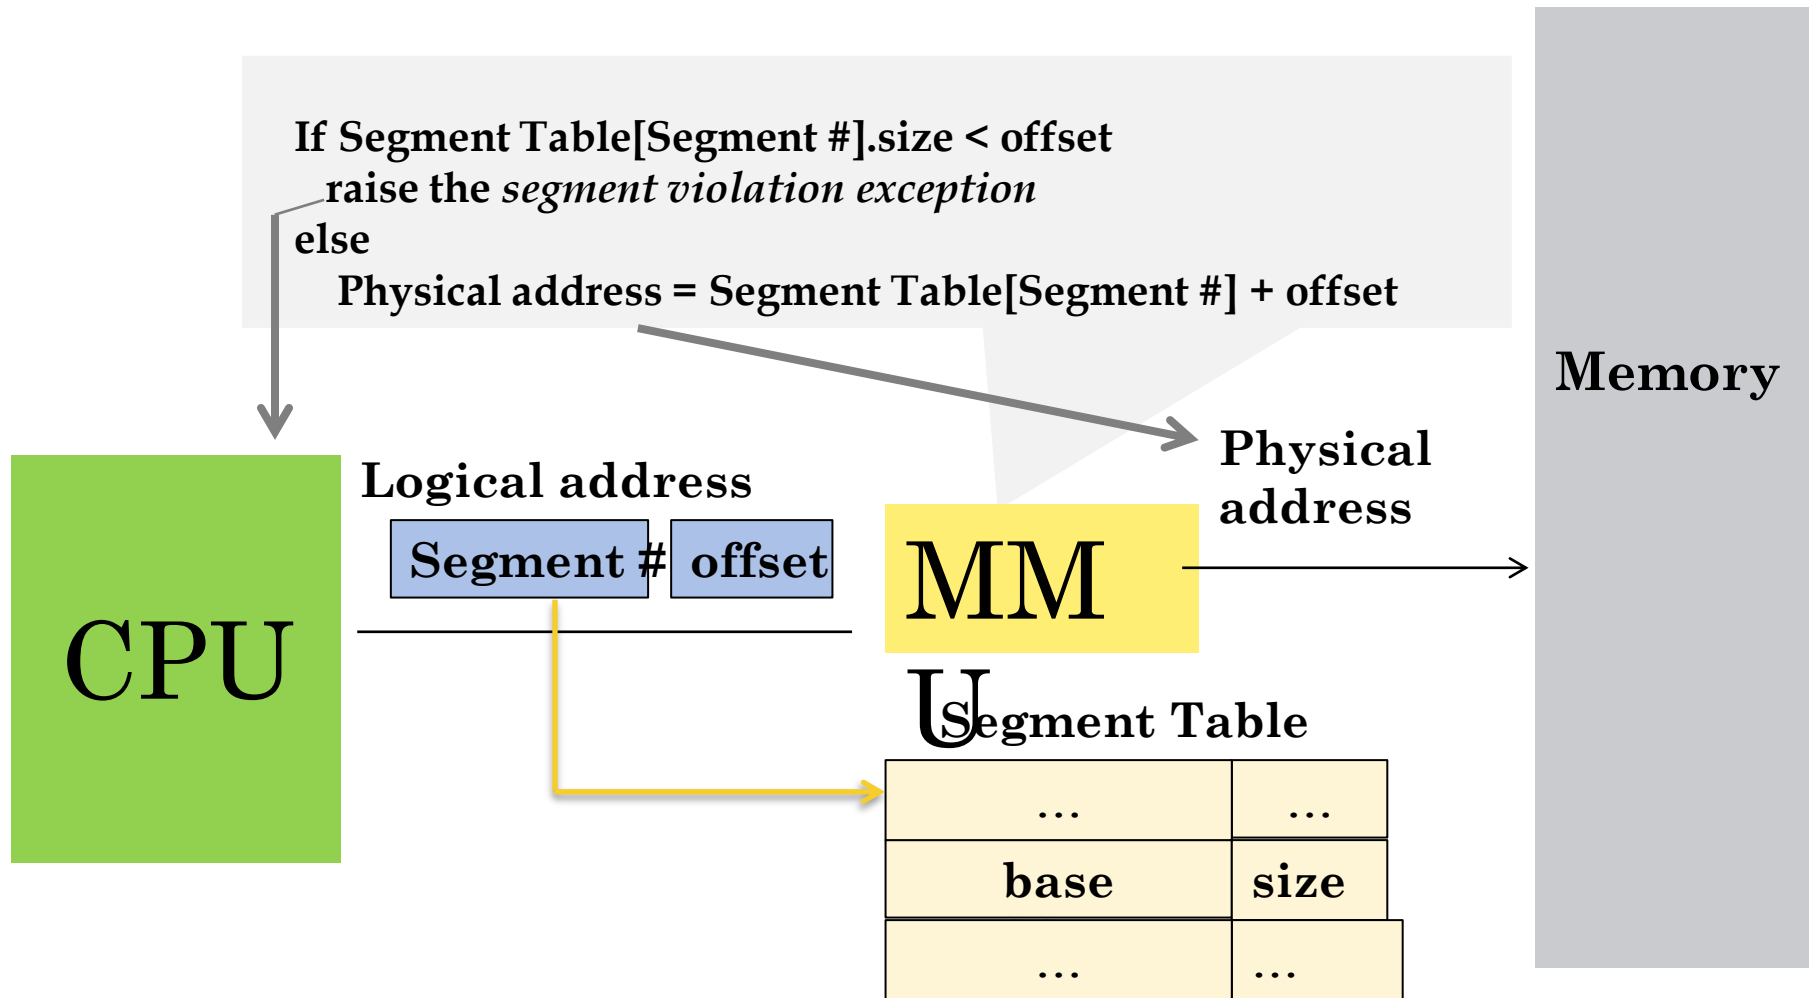
Memory Management Unit (MMU) translates a virtual address into a “real” (physical) address




MAR: Memory Address Register

MDR: Memory Data Register

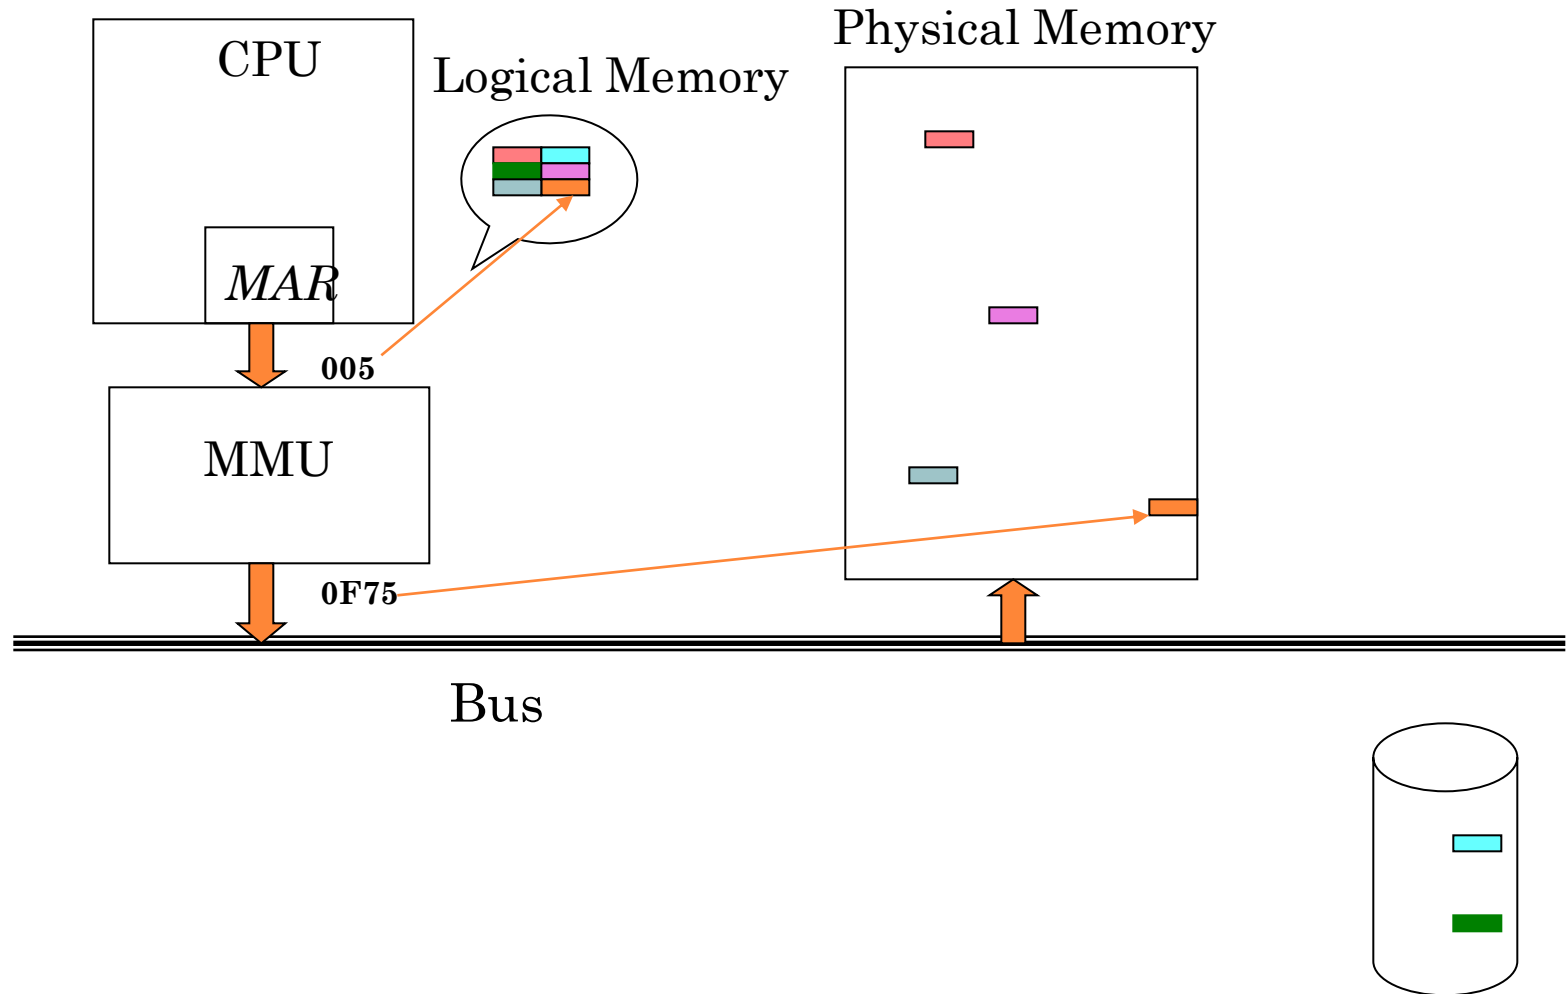
Segmentation: The MMU translation processing



MEMORY REFERENCE TRANSLATION WITH THE *MEMORY MANAGEMENT UNIT (MMU)*



VIRTUAL MEMORY: CONTIGUOUS AND “INFINITE”



PAGES: VIRTUAL MEMORY, PHYSICAL MEMORY, DISC

0	A ₁
1	A ₂
2	A ₃
3	A ₄
4	A ₅
5	A ₆

The Process A
Virtual Memory

0	B ₁
1	B ₂
2	B ₃
3	B ₄
4	B ₅
5	B ₆

The Process B
Virtual Memory

Free Frame List:
6,7

0	A ₁
1	B ₄
2	A ₆
3	A ₂
4	B ₃
5	B ₅
6	x
7	x

Physical
Memory


A ₁	A ₂	A ₃	A ₄
A ₅	A ₆	B ₁	B ₂
B ₃	B ₄	B ₅	B ₆

Disk

ACCESS PROTECTION RINGS

Segment Table revisited


...
Start	size	ring	Access type = {read, write, execute}
...



Code permitted to access ring i may access ring j if and only if $j > i$,

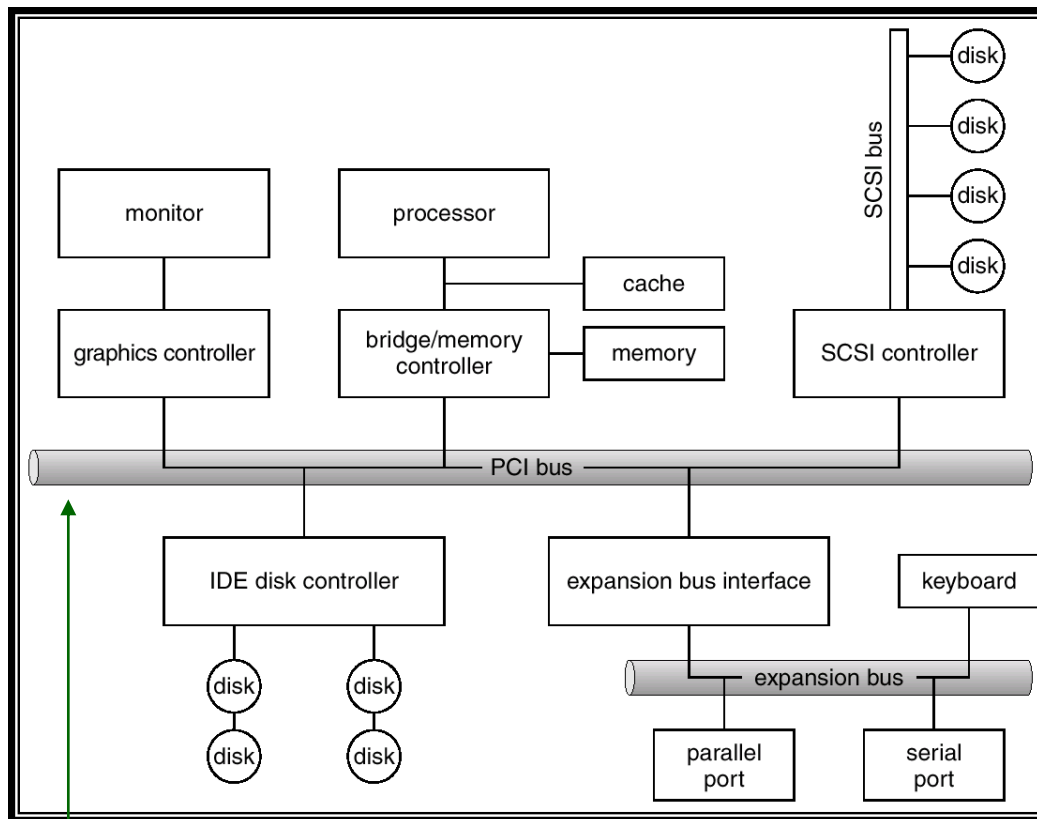
ANOTHER BIT OF HISTORY

- *Multics* had **eight** rings
- *OS/2* uses **three** rings: ring 0 for kernel code and device drivers; ring 2 for privileged code (user programs with I/O access permissions); and ring 3, for unprivileged access
- *UNIX and Microsoft Windows NT* support rings 0 (kernel) and 2 (user programs)



Conclusion: An OS kernel by itself cannot create an isolated virtual machine

A PHYSICAL MACHINE



Note: *Integrated Disk Electronics (IDE)* is the least expensive current disk technology. IDE support is usually built into the main board.

Small Computer Systems Interface (SCSI) supports disks, tapes, and CD-ROM drives. While IDE disks provide up to one gigabyte of storage, SCSI disks are available with four to 9 gigabytes of storage.

Peripheral Component Interconnection (PCI)

Igor Faynberg, Cloud Computing, Lecture 2

SUMMARY

- A process as a unit has its own virtual world of resources: A CPU, all I/O devices, an “infinite” memory
- A process relies on an operating system—the *government*, which ensures that all processes are treated fairly
- The operating system *kernel* is the only entity that can execute privileged instructions; the kernel is entered via *traps* (internally) and *interrupts* caused by other events
- A process is aware of other processes
- A *machine* (Hardware + OS) executes many processes
- Question: How can we *create* a virtual machine?